

Garima Diyawar

ndiyawar@gmail.com | linkedin.com/in/garima-diyawar | github.com/garimadiyawar

TECHNICAL SKILLS

LLM and GenAI: LangChain, LangGraph, RAG (Retrieval-Augmented Generation), Multi-agent Orchestration, Prompt Engineering, Agent Tracing (LangSmith), Hallucination Detection, Ollama (OpenAI-compatible API), Hybrid Retrieval (BM25 + Vector)

Vector Search and Databases: ChromaDB, Embeddings, Semantic Search, Azure Cognitive Search, TinyDB

NLP and Transformers: Hugging Face Transformers, RoBERTa Fine-tuning, Multi-label Text Classification, Named Entity Recognition, Sentiment Analysis (Azure Text Analytics)

APIs and Services: OpenFDA API, Azure Text Analytics, Azure Translator API, Azure Cognitive Services

Model Serving: FastAPI, REST API, Streamlit, Flask, GitHub Actions CI/CD

Languages: Python, SQL, Bash

Data and ML: Pandas, NumPy, scikit-learn, XGBoost, LightGBM, Matplotlib

EXPERIENCE

AI on Azure Intern | Edunet Foundation | Microsoft Initiative | AICTE

May 2025 - Jun 2025 | Remote

- Built a knowledge mining and compliance risk detection pipeline using Azure Cognitive Search with AI enrichment skillsets, Azure Text Analytics (NER and sentiment), and Azure Translator API to process legal and financial documents.
- Designed a retrieval and classification workflow surfacing entity-linked compliance signals; deployed as a live Streamlit app with GitHub Actions CI/CD at github.com/garimadiyawar/Knowledge-Mining-and-Compliance-Risk-Detection-System.

PROJECTS

Enterprise Brain - Multi-Agent RAG System | Python, LangGraph, LangChain, ChromaDB, Ollama, Streamlit

- Built a multi-agent RAG pipeline that ingests PDFs, Slack messages, and emails and answers natural-language queries with source citations; used LangGraph to orchestrate retrieval, reasoning, and answer-verification agents.
- Implemented hybrid BM25 and vector retrieval with ChromaDB and an Ollama-hosted LLM (OpenAI-compatible API); automated fact-check agent validates responses to minimize hallucinations.
- Added role-based access control (4 user roles), citation tracing, and a live evaluation dashboard; fully local with no external API costs.

Pharma Adverse Event Monitor - Agentic Safety Pipeline | Python, LangGraph, LangChain, OpenFDA API, LangSmith, Streamlit

- Designed a 3-agent LangGraph pipeline: a retrieval agent fetches adverse event data from the OpenFDA FAERS API, a reasoning agent classifies signals (RED/YELLOW/GREEN) by organ system, and a reporting agent outputs structured safety summaries.
- Integrated LangSmith for end-to-end agent tracing and prompt debugging; built an interactive Streamlit dashboard for browsing reports and monitoring agent steps in real time.

BeNice - Implicit Bias Detector (Transformer Fine-tuning) | Python, PyTorch, RoBERTa, Hugging Face Transformers, scikit-learn

- Fine-tuned RoBERTa-base for multi-label classification of online harm (misogyny, racism, casteism, cyberbullying) using BCEWithLogitsLoss with class weighting; achieved 75% micro F1 and 41% eval-loss reduction over 4 epochs.
- Applied per-label threshold optimization for independent precision/recall control per class; annotated 200+ real-world samples and documented model failure modes on sarcasm and reclaimed language.

EDUCATION

B.Tech, Computer Science | Kalinga University, Raipur

2022 - 2026

CGPA: 8.82/10 | Aggregate: 84.24% | Class X: 90.8% | Class XII: 80.4%

CERTIFICATIONS

- Google Advanced Data Analytics Specialization - Coursera/Google
- Google Data Analytics Professional Certificate - Coursera/Google
- Google IT Automation with Python Professional Certificate - Coursera/Google